

# Capability-Based Memory Protection for Scalable Vector Processing

Samuel Stark ([sws35@cam.ac.uk](mailto:sws35@cam.ac.uk))

June 16th 2022

## Capability-Based Memory Protection for Scalable Vector Processing

## Capability-Based Memory Protection for Scalable Vector Processing

## Capability-Based Memory Protection for Scalable Vector Processing

### Capability-Based Memory Protection == CHERI<sup>1</sup>

- Memory is normally addressed with integers
  - Integer addresses can be *forged*
  - Code can be *tricked* into accessing memory it shouldn't
- CHERI architectures use *capabilities* instead
  - Capability = Bounds + current address
- Capabilities cannot be forged, only *derived* from other capabilities

Code can only access memory when it has been *given* access to that memory

<sup>1</sup>Capability Hardware Enhanced RISC Instructions[1]

### (Scalable) Vector Processing

- Vector architectures allow programmers to use SIMD
- Most vector architectures have *fixed-length* vectors
  - SSE, Arm Neon = 128-bit, AVX-512 = 512
  - Vector lengths have hardware tradeoff
  - Need to recompile code for different vector lengths
- *Scalable vector* architectures give designers flexibility[2]
  - Code doesn't rely on fixed vector length

24	9	9	28
+			
13	17	10	24
=			
37	26	19	52

**Table 1:** Vector addition

## Have we combined them before?

- CHERI affects vector memory accesses
  - Loading N vector elements in a single instruction
  - Per-element bounds checks could be expensive?
- Arm have manufactured CHERI hardware
  - Has fixed-length SIMD
  - Doesn't support complex access patterns
- No other general-purpose CHERI processors with vector support

Where does this matter?

### memcpy!


- Take data and copy it somewhere else
- Extremely widespread operation
- Vectors can copy more data per instruction
- Vectorized **memcpy** should work and go fast on CHERI





### Vectorized memcpy!

- Take data and copy it somewhere else
- Extremely widespread operation
- Vectors can copy more data per instruction
- Vectorized **memcpy** should work and go fast on CHERI



Make vectorized `memcpy` functional and fast on CHERI

Combine the RISC-V Vector extension (RVV) with CHERI-RISC-V

1. Write a RISC-V CHERI-RVV emulator in Rust
  - Demonstrates hardware feasibility
2. Write test programs in C
  - Demonstrates software feasibility
3. Run the test programs on the emulator!

Make vectorized `memcpy` functional and fast on CHERI

Combine the RISC-V Vector extension (RVV) with CHERI-RISC-V

## RVV (original)

- Uses integer addressing
- Loads/stores integer data

## CHERI-RVV (ours)

- Uses **capability** addressing
  - Performance concerns?
- Loads/stores integers **and capabilities**
- Doesn't break CHERI security

Step 1: Making vector accesses *use*  
capabilities



# Vector memory access patterns

## 1. Unit-stride

- $\text{base}$ ,  $\text{base}+1$ ,  
 $\text{base}+2 \dots$

## 2. Strided

- $\text{base}$ ,  $\text{base}+n$ ,  
 $\text{base}+2n \dots$

## 3. Indexed

- $\text{base} + \text{offset}[0]$ ,  
 $\text{base} + \text{offset}[1]$ ,  
 $\text{base} + \text{offset}[2] \dots$




Figure 1: Unit access

# Vector memory access patterns

## 1. Unit-stride

- $\text{base}$ ,  $\text{base}+1$ ,  
 $\text{base}+2\dots$

## 2. Strided

- $\text{base}$ ,  $\text{base}+n$ ,  
 $\text{base}+2n\dots$

## 3. Indexed

- $\text{base} + \text{offset}[0]$ ,  
 $\text{base} + \text{offset}[1]$ ,  
 $\text{base} + \text{offset}[2]\dots$





Figure 2: Strided

- base, base+1, base+2...

- base, base+n, base+2n...

- `base + offset[0],`  
`base + offset[1],`  
`base + offset[2]...`



### Figure 3: Indexed

# Vector memory access patterns

## Unit-stride

Base address

0x37f0



## Strided

Base address

0x37f0

Stride

2




## Indexed

Base address

0x37f0

Index vector

8	2	13	4
---	---	----	---





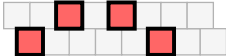
Property

Example value

Memory pattern



# Vector memory access patterns

Unit-stride					
Base <span>capability</span>	0x3700.. 0x37f0 ..0x3800				
Strided					
Base <span>capability</span>	0x3700.. 0x37f0 ..0x3800				
Stride	2				
Indexed					
Base <span>capability</span>	0x3700.. 0x37f0 ..0x3800				
Index vector	<table><tr><td>8</td><td>2</td><td>13</td><td>4</td></tr></table>		8	2	13
8	2	13	4		
Property	Example value	Memory pattern			

## Step 2: Making vector accesses *copy* capabilities




# Storing capabilities in memory

- Memory can hold both capabilities and integers
- Separate *tag memory* denotes which regions are capabilities
  - Access to tag memory is controlled by hardware
- Without the tag, you get the *integer encoding* of the capability

tag	data[128:0]
int	4436c97773d3504f
cap	0x3700 .. 0x37f0 .. 0x3800
...	
int	

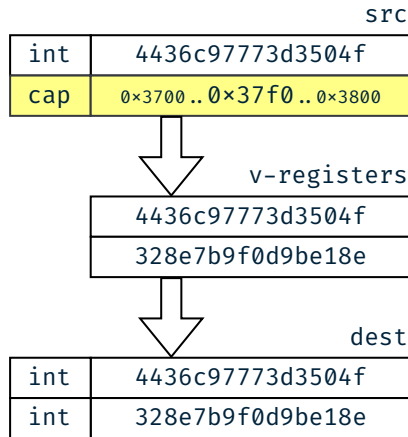
# Storing capabilities in memory

- Memory can hold both capabilities and integers
- Separate *tag memory* denotes which regions are capabilities
  - Access to tag memory is controlled by hardware
- Without the tag, you get the *integer encoding* of the capability



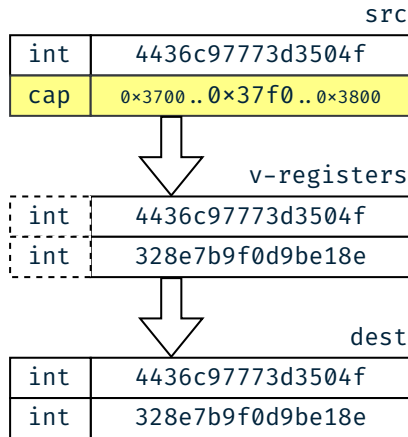
# Integer-only memcpy

- The original RVV specification doesn't consider capabilities
- Assumes vectors only hold **integer** data
- $\Rightarrow$  **memcpy** converts capabilities to integers  
: (



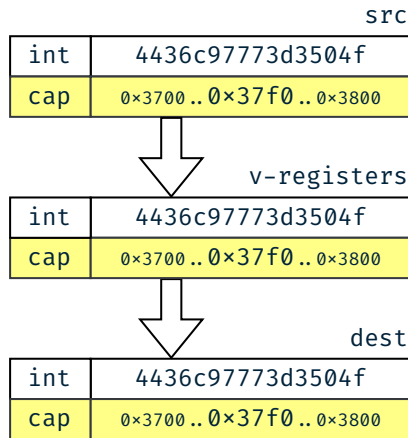
# Integer-only memcpy

- The original RVV specification doesn't consider capabilities
- Assumes vectors only hold **integer** data
- $\Rightarrow$  **memcpy** converts capabilities to integers  
: (



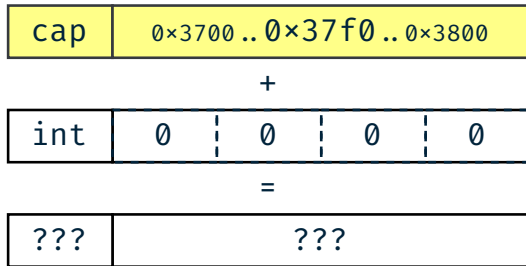
# Copying capabilities in vectors

- If we add tag bits to vector registers, we can load/store them correctly
- But does that make anything else more complicated?
  - Yes



## Storing capabilities in vectors???

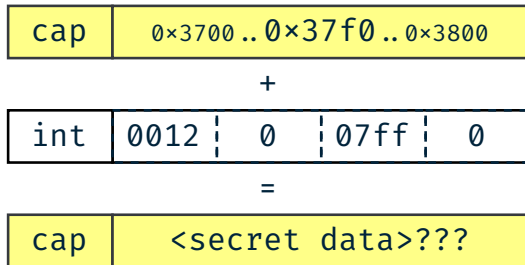
- Now *all* vector instructions can interact with capabilities!
- If we aren't careful, attackers could forge capabilities
- We introduce two contexts of accessing vector registers
  - Integer context
  - Capability context






## Storing capabilities in vectors???

- Now *all* vector instructions can interact with capabilities!
- If we aren't careful, attackers could forge capabilities
- We introduce two contexts of accessing vector registers
  - Integer context
  - Capability context

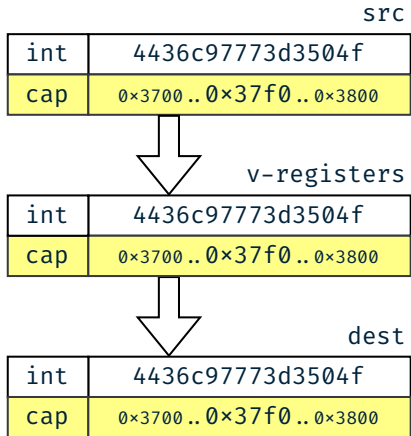


## Storing capabilities in vectors???


- Now *all* vector instructions can interact with capabilities!
- If we aren't careful, attackers could forge capabilities
- We introduce two contexts of accessing vector registers
  - Integer context
  - Capability context



# Integer/Capability context



Capability context  
(128-bit vector loads/stores)



Integer context  
(Everything else)

## memcpy works!

	RV32	RV-64				
	llvm-13	llvm-13	llvm-15	gcc	CHERI	CHERI (Int)
Copy	Y	Y	Y	Y	Y	Y
Copy + Invalidate	-	-	-	-	Y	Y

## CHERI-RVV Summary

---

Uses capability addressing

Loads/stores integers and capabilities

Doesn't break CHERI security

---

Supports all vanilla RVV instructions

Is binary-compatible with vanilla RVV

Can be\* source-compatible with vanilla RVV

---

Has a reference implementation:

Emulator, compiler\*, test programs

9,500 LoC

---

\*compiler requires engineering work

## Future work

Do more with vectors than just

1. Vectorized **memcpy**
2. Vectorized tag clearing

Add new instructions for e.g.  
temporal revocation[3]?

Samuel Stark  
[sws35@cam.ac.uk](mailto:sws35@cam.ac.uk)

## CHERI-RVV Summary

---

Uses capability addressing

Loads/stores integers and capabilities

Doesn't break CHERI security

---

Supports all vanilla RVV instructions

Is binary-compatible with vanilla RVV

Can be\* source-compatible with vanilla RVV

---

Has a reference implementation:

Emulator, compiler\*, test programs

9,500 LoC

---

\*compiler requires engineering work

## Future work

Do more with vectors than just

1. Vectorized `memcpy`
2. Vectorized tag clearing

Add new instructions for e.g.  
temporal revocation[3]?

Samuel Stark  
[sws35@cam.ac.uk](mailto:sws35@cam.ac.uk)

## Per-element checks

- Vector hardware can coalesce element accesses
  - e.g. 4x 32-bit elements in the same cache line can be transferred over a 128-bit bus at once
- Want to coalesce the per-element capability checks as well
  - Otherwise they could bottleneck
  - Or use too much logic
- We found we *can* coalesce capability checks if they succeed
  - i.e. "is the cache line inside the capability bounds"
  - But if that check fails, we have to check each element individually
  - RVV requires that any synchronous exception (i.e. capability check) reports the element that triggered it

### References

---

- [1] Robert N M Watson et al. *An Introduction to CHERI*. UCAM-CL-TR-941. September 2019, p. 43.
- [2] Nigel Stephens et al. “The ARM Scalable Vector Extension”. In: *IEEE Micro* 37.2 (March 2017), pp. 26–39. ISSN: 0272-1732. DOI: [10.1109/MM.2017.35](https://doi.org/10.1109/MM.2017.35).
- [3] Hongyan Xia et al. “CHERIvoke: Characterising Pointer Revocation Using CHERI Capabilities for Temporal Memory Safety”. In: (2019), p. 14. DOI: [10/gm9ngg](https://doi.org/10/gm9ngg).